

Time Series

```
import pandas as pd
import matplotlib.pyplot as plt
from statsmodels.tsa.stattools import adfuller
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.tsa.arima.model import ARIMA
import seaborn as sns
import numpy as np
```

loading and viewing data

```
temp_df = pd.read_csv("C:\my files\DailyDelhiClimateTrain.csv")
```

```
temp_df.head()
```

	date	meantemp	humidity	wind_speed	meanpressure
0	2013-01-01	10.000000	84.500000	0.000000	1015.666667
1	2013-01-02	7.400000	92.000000	2.980000	1017.800000
2	2013-01-03	7.166667	87.000000	4.633333	1018.666667
3	2013-01-04	8.666667	71.333333	1.233333	1017.166667
4	2013-01-05	6.000000	86.833333	3.700000	1016.500000

```
temp_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1462 entries, 0 to 1461
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   date            1462 non-null   object
1   meantemp        1462 non-null   float64
2   humidity        1462 non-null   float64
3   wind_speed      1462 non-null   float64
4   meanpressure    1462 non-null   float64
dtypes: float64(4), object(1)
memory usage: 57.2+ KB
```

setting date as index

```
#checking for nulls in date column
print(temp_df[temp_df["date"].isna()])
```

```
Empty DataFrame
```

```
Columns: [date, meantemp, humidity, wind_speed, meanpressure]
```

```
Index: []
```

```
#convert object into date datatype
temp_df["date"] = pd.to_datetime(temp_df["date"],errors="coerce")
temp_df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1462 entries, 0 to 1461
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   date             1462 non-null   datetime64[ns]
1   meantemp         1462 non-null   float64
2   humidity         1462 non-null   float64
3   wind_speed       1462 non-null   float64
4   meanpressure     1462 non-null   float64
dtypes: datetime64[ns](1), float64(4)
memory usage: 57.2 KB
```

```
#setting index
temp_df.set_index("date",inplace = True)
temp_df.head()
```

	meantemp	humidity	wind_speed	meanpressure
date				
2013-01-01	10.000000	84.500000	0.000000	1015.666667
2013-01-02	7.400000	92.000000	2.980000	1017.800000
2013-01-03	7.166667	87.000000	4.633333	1018.666667
2013-01-04	8.666667	71.333333	1.233333	1017.166667
2013-01-05	6.000000	86.833333	3.700000	1016.500000

visualise

```
temp_df.plot(figsize=(12,6),subplots = True)
plt.show()
```



conclusion: 1) meantemp - every year it goes on increases and decreases - when ever the humidity is high the meantemp decreases 2) humidity - humidity varies slighty 3) wind_speed - it is high btw 2013 and 2014 - it is normal in 2016 to 2017 4) meanpressure - it maintance same range up to 2016 from 2013 - it is high in 2016 to 2017

stationarity

A time series is stationary if its statistical properties(men,variance,autocorrelation)remain constant over time

hypotheses of the ADF Test

- Null Hypothesis:The time series has a unit root (non - stationary)
- Alternative Hypothesis:The time serie does **not** have a unit root(stationary)

```
adfuller_result = adfuller(temp_df["meantemp"])
print(adfuller_result)
```

```
(-2.0210690559206728, 0.27741213723016056, 10, 1451, {'1%': -
3.4348647527922824, '5%': -2.863533960720434, '10%': -
2.567831568508802}, 5423.895746470953)
```

```
if adfuller_result[1]<0.05:
    print("Stationary")
else:
    print("Non - Stationary")
```

Non - Stationary

Differencing to Remove Trend

- Differencing is a technique used to make a non-stationary time series stationary by removing trends or seasonality. It involves subtracting the previous observation from the current observation
- `temp=[20,21,22,24,25,27,28,27]`
- `diff=[1,1,2,1,2,1,-1]`
- the new series fluctuates around zero -2 to 2


```
temp_df['meantemp_diff'] = temp_df["meantemp"].diff()  
temp_df.head()
```

	meantemp	humidity	wind_speed	meanpressure
meantemp_diff				
date				
2013-01-01	10.000000	84.500000	0.000000	1015.666667
NaN				
2013-01-02	7.400000	92.000000	2.980000	1017.800000
2.600000				-
2013-01-03	7.166667	87.000000	4.633333	1018.666667
0.233333				-
2013-01-04	8.666667	71.333333	1.233333	1017.166667
1.500000				
2013-01-05	6.000000	86.833333	3.700000	1016.500000
2.666667				-

```
adfuller_resul_afterdiff = adfuller(temp_df["meantemp_diff"].dropna())  
if adfuller_result[1]>0.05:  
    print("Stationary")  
else:  
    print("Non - Stationary")
```

Stationary

```
temp_df["meantemp_diff"].plot(title = "Difference Mean Temperature")  
plt.show()
```




conclusion: 1) The data shows variation in temperature between -10 and 7.5 with a mean of 0 2) there is no long difference in trends but frequent peaks shows low spikes in 2015-2016. 3) there is outliers 4) the mean temperature is now stationary

use seasonal decomposition to analyze trend, seasonality, and residuals

- Trend : The long term pattern(increase or decrease over time)
- seasonality : The repeating patterns at fixed intervals(monthly sales spikes)
- Residual(Noise) : The random variations that are not explained by trend or seasonality
- interpreting the Output:
- original series : The raw time series data
- trend component : The general direction of the data over time
- seasonal component : The repeating patterns(eg: higher sales in december)
- Residual component : The remaining part after removing trend and seasonality

```
decomposing = seasonal_decompose(temp_df["meantemp"], model = "additive", period = 365)
```

```
decomposing.plot()
plt.show()
```



conclusion: 1) there is a rapid increase in trend between 7/2015 to 7/2016 2) it exhibits a clear seasonal pattern, with same fluctuations repeating over time 3) the residuals are randomly distributed

ARIMA

```
# splitting data
len(temp_df)

1462

print(len(temp_df)*0.8)

1169.6000000000001

train = temp_df.iloc[0:1169]
test = temp_df.iloc[1169:]
len(test)
```

293

```
mymodel = ARIMA(train["meantemp"],order=(1,1,1))
```

```
C:\ProgramData\anaconda3\Lib\site-packages\statsmodels\tsa\base\
tsa_model.py:473: ValueWarning: No frequency information was provided,
so inferred frequency D will be used.
```

```
self._init_dates(dates, freq)
```

```
C:\ProgramData\anaconda3\Lib\site-packages\statsmodels\tsa\base\
tsa_model.py:473: ValueWarning: No frequency information was provided,
so inferred frequency D will be used.
```

```
self._init_dates(dates, freq)
```

```
C:\ProgramData\anaconda3\Lib\site-packages\statsmodels\tsa\base\
tsa_model.py:473: ValueWarning: No frequency information was provided,
so inferred frequency D will be used.
```

```
self._init_dates(dates, freq)
```

```
mymodel = mymodel.fit()
```

```
forecast = mymodel.forecast(steps = len(test))
```

```
print(forecast.head())
```

```
2016-03-15    22.826205
```

```
2016-03-16    23.085687
```

```
2016-03-17    23.234913
```

```
2016-03-18    23.320731
```

```
2016-03-19    23.370084
```

```
Freq: D, Name: predicted_mean, dtype: float64
```

```
test["forecaste"] = forecast
```

```
test.head()
```

```
C:\Users\DELL\AppData\Local\Temp\ipykernel_9892\3744547450.py:1:
```

```
SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation:
```

```
https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#
returning-a-view-versus-a-copy
```

```
test["forecaste"] = forecast
```

```
          meantemp  humidity  wind_speed  meanpressure
meantemp_diff \
date
```

```
2016-03-15    24.066667   58.933333     8.646667    1014.866667
```

```
1.691667
```

```
2016-03-16    23.937500   53.750000    10.881250    1012.812500
```

```
0.129167
```

```
2016-03-17    26.312500   50.312500     6.843750    1010.437500
```

```

2.375000
2016-03-18  26.187500  61.250000    6.712500   1009.812500   -
0.125000
2016-03-19  26.785714  61.857143    3.578571   1009.214286
0.598214

```

```


          forecaste
date
2016-03-15  22.826205
2016-03-16  23.085687
2016-03-17  23.234913
2016-03-18  23.320731
2016-03-19  23.370084

```

```

plt.figure(figsize=(8,5))
plt.plot(test.index,test["meantemp"],color="orange",label="Original")
plt.plot(test.index,test["forecaste"],color="green",label="Forecast")
plt.title("ORIGINAL VS FORECAST")
plt.show()

```



```
len(temp_df)
```

```
1462
```



```

print(len(temp_df)*0.8)

1169.6000000000001

train1 = temp_df.iloc[0:1169]
test1 = temp_df.iloc[1169:]
len(test1)

293

mymodel = ARIMA(train["meantemp_diff"],order=(1,1,1))

C:\ProgramData\anaconda3\Lib\site-packages\statsmodels\tsa\base\
tsa_model.py:473: ValueWarning: No frequency information was provided,
so inferred frequency D will be used.
    self._init_dates(dates, freq)
C:\ProgramData\anaconda3\Lib\site-packages\statsmodels\tsa\base\
tsa_model.py:473: ValueWarning: No frequency information was provided,
so inferred frequency D will be used.
    self._init_dates(dates, freq)
C:\ProgramData\anaconda3\Lib\site-packages\statsmodels\tsa\base\
tsa_model.py:473: ValueWarning: No frequency information was provided,
so inferred frequency D will be used.
    self._init_dates(dates, freq)

mymodel = mymodel.fit()

forecast = mymodel.forecast(steps = len(test))
print(forecast.head())

2016-03-15    0.322914
2016-03-16   -0.040400
2016-03-17    0.019656
2016-03-18    0.009729
2016-03-19    0.011370
Freq: D, Name: predicted_mean, dtype: float64

test1["forecaste_diff"] = forecast
test1.head()

C:\Users\DELL\AppData\Local\Temp\ipykernel_9892\990163927.py:1:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#
returning-a-view-versus-a-copy
    test1["forecaste_diff"] = forecast

      meantemp  humidity  wind_speed  meanpressure
meantemp_diff \

```

date


2016-03-15	24.066667	58.933333	8.646667	1014.866667	
2016-03-16	23.937500	53.750000	10.881250	1012.812500	-
2016-03-17	26.312500	50.312500	6.843750	1010.437500	
2016-03-18	26.187500	61.250000	6.712500	1009.812500	-
2016-03-19	26.785714	61.857143	3.578571	1009.214286	

forecaste_diff

date

2016-03-15	0.322914
2016-03-16	-0.040400
2016-03-17	0.019656
2016-03-18	0.009729
2016-03-19	0.011370

```
plt.figure(figsize=(8,5))
plt.plot(test.index,test["meantemp_diff"],color="orange",label="Original")
plt.plot(test.index,test["forecaste_diff"],color="green",label="Forecast")
plt.title("ORIGINAL VS FORECAST")
plt.show()
```



conclusion: 1) 15th mar, original value = 24.06667 , model says there is -0.040400 change on next day $24.06667 - 0.040400 = 24.10$ (predict) and 23.937(original)